

TP 4 – AOS1

PCA

1 Python warm up: PCA by hand

- ① Generate a dataset with the following instruction

```
| X = np.random.multivariate_normal([1, 3], [[2, 1], [1, 2]], 100)
```

How many samples are generated? How many features? What is the underlying distribution of samples in X ?

- ② Verify the relation that exists between singular values and eigenvalues using a matrix X . To use the functions provided by the `scipy` library, use the following command:

```
| import scipy.linalg as linalg
```

and look at the functions `linalg.eig`, `linalg.eigh`, `linalg.eigvals`, `linalg.eigvalsh`, `linalg.svd` `linalg.svdvals`

- ③ Compute the principal directions and principal components by hand using the unbiased variance-covariance estimator. Verify that they coincide with the ones computed by `scikit-learn`.

2 PCA for dimension reduction

In this section, we use the `house_prices` regression dataset. To load it use

```
| from sklearn.datasets import fetch_openml
| housing = fetch_openml(name="house_prices", as_frame=True)
| X = housing.frame
| X = X.loc[:, ~X.isna().any() & (X.dtypes.astype(str).isin(["float64", "int64"]))]
| y = housing.target
```

- ④ Perform a PCA on this dataset and study how many number of principal components should be retained from the two empirical methods seen in class.
- ⑤ Describe the following code. What is it supposed to be doing? Adapt it to determine the optimal number of principal components for the regression task at hand.

```

from sklearn.decomposition import PCA
from sklearn.linear_model import LinearRegression
from sklearn.model_selection import GridSearchCV
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import StandardScaler

pca = PCA()
lin = LinearRegression()
pca_lin = Pipeline([("scale", StandardScaler()), ("pca", pca), ("lin", lin)])
clf = GridSearchCV(
    estimator=pca_lin,
    cv=10,
    param_grid=dict(pca__n_components=range(1, X.shape[1] + 1)),
)
clf.fit(X, y)

```

3 Problem: band reduction in multispectral images

A multispectral image is an image that has several components. For example, a color image has 3 components: red, green and blue and each pixel can be viewed as a vector in \mathbb{R}^3 . More generally a multispectral image of size $N \times M$ with P spectral bands can be stored as a $N \times M \times P$ array. There are $N \times M$ pixels living in \mathbb{R}^P .

When the number of spectral bands P is too large, it is desirable to somehow reduce that number ultimately to 3 for viewing purposes. This process is called band reduction.

Propose a method using the PCA performing a band reduction to 3 bands and use it on the provided multispectral image.

Some multispectral images are available on the internet to test your band reduction algorithm. See for example the following website

- <http://lesun.weebly.com/hyperspectral-data-set.html>

Most of them are available as a Matlab data file (.mat files). It can be loaded with `scipy` with the following function

```
| scipy.io.loadmat
```

You will probably have to reshape arrays. It can be done with the `reshape` method. For example, an array of size $6 \times 6 \times 3$ can be “linearized” using `reshape`

```
| X_lin = X.reshape((-1, 3))
```

the `-1` is automatically inferred from the number of elements in the array. The array is then reshaped into an array of size 36×3 .

It might be handy to be able to rescale the data when it has to belong to some specific range. `scikit-learn` has several rescalers available. For example

```
| from sklearn.preprocessing import MinMaxScaler
```

rescales the data between 0 and 1.

`matplotlib` can display images with the function

| `plt.imshow`

Beware of the type of the array (float or integers)!