

TP 4 – AOS1

PCA

Corrigé

1 Python warm up: PCA by hand

- ① Generate a dataset with the following instruction

```
| X = np.random.multivariate_normal([1, 3], [[2, 1], [1, 2]], 100)
```

How many samples are generated? How many features? What is the underlying distribution of samples in X ?

```
In [1]: | import matplotlib.pyplot as plt
         | import numpy as np
         | import scipy.linalg as linalg
         | X = np.random.multivariate_normal([1, 3], [[2, 1], [1, 2]], 100)
         | X.shape
```

```
Out [1]: | (100, 2)
```

There are 100 samples and 2 features. Samples are drawn according to $\mathcal{N}\left(\begin{pmatrix} 1 \\ 3 \end{pmatrix}, \begin{pmatrix} 2 & 1 \\ 1 & 2 \end{pmatrix}\right)$.

- ② Verify the relation that exists between singular values and eigenvalues using a matrix X . To use the functions provided by the `scipy` library, use the following command:

```
| import scipy.linalg as linalg
```

and look at the functions `linalg.eig`, `linalg.eigh`, `linalg.eigvals`, `linalg.eigvalsh`, `linalg.svd` `linalg.svdvals`

```
In [2]: | X = np.random.normal(size=(6, 2))
         | print(linalg.eigvalsh(X.T @ X))
```

```
Out [2]: | [0.64178977 7.47512647]
```

```
In [3]: | print(linalg.svdvals(X)**2)
```

```
Out [3]: | [7.47512647 0.64178977]
```

Nonzero eigenvalues of $X^T X$ (or XX^T) are squared singular values of X .

- ③ Compute the principal directions and principal components by hand using the unbiased variance–covariance estimator. Verify that they coincide with the ones computed by `scikit-learn`.

```
In [4]: n, p = 100, 15
        X = np.random.normal(size=(n, p))
        X0 = X - X.mean(axis=0)
        V = 1/(n-1) * X0.T @ X0
        vp, U = linalg.eigh(V)
        print(vp)

Out [4]: [0.39346284 0.46841904 0.54507063 0.56325593 0.72062481
          ↪ 0.78908287
          0.87556462 0.92644902 0.96237418 1.15910771 1.18983685
          ↪ 1.34312761
          1.37139569 1.62307881 1.82253541]

In [5]: Xpca = X0 @ U
        print(n / (n-1) * Xpca.std(axis=0)**2)

Out [5]: [0.39346284 0.46841904 0.54507063 0.56325593 0.72062481
          ↪ 0.78908287
          0.87556462 0.92644902 0.96237418 1.15910771 1.18983685
          ↪ 1.34312761
          1.37139569 1.62307881 1.82253541]

In [6]: from sklearn.decomposition import PCA
        pca = PCA()
        pca.fit(X)

Out [6]: PCA()

In [7]: print(pca.explained_variance_)

Out [7]: [1.82253541 1.62307881 1.37139569 1.34312761 1.18983685
          ↪ 1.15910771
          0.96237418 0.92644902 0.87556462 0.78908287 0.72062481
          ↪ 0.56325593
          0.54507063 0.46841904 0.39346284]
```

We have the same eigenvalues (in reverse order).

2 PCA for dimension reduction

In this section, we use the `house_prices` regression dataset. To load it use

```
from sklearn.datasets import fetch_openml
housing = fetch_openml(name="house_prices", as_frame=True)
X = housing.frame
X = X.loc[:, ~X.isna().any() & (X.dtypes.astype(str).isin(["float64", "int64"]))]
y = housing.target
```

- ④ Perform a PCA on this dataset and study how many number of principal components should be retained from the two empirical methods seen in class.

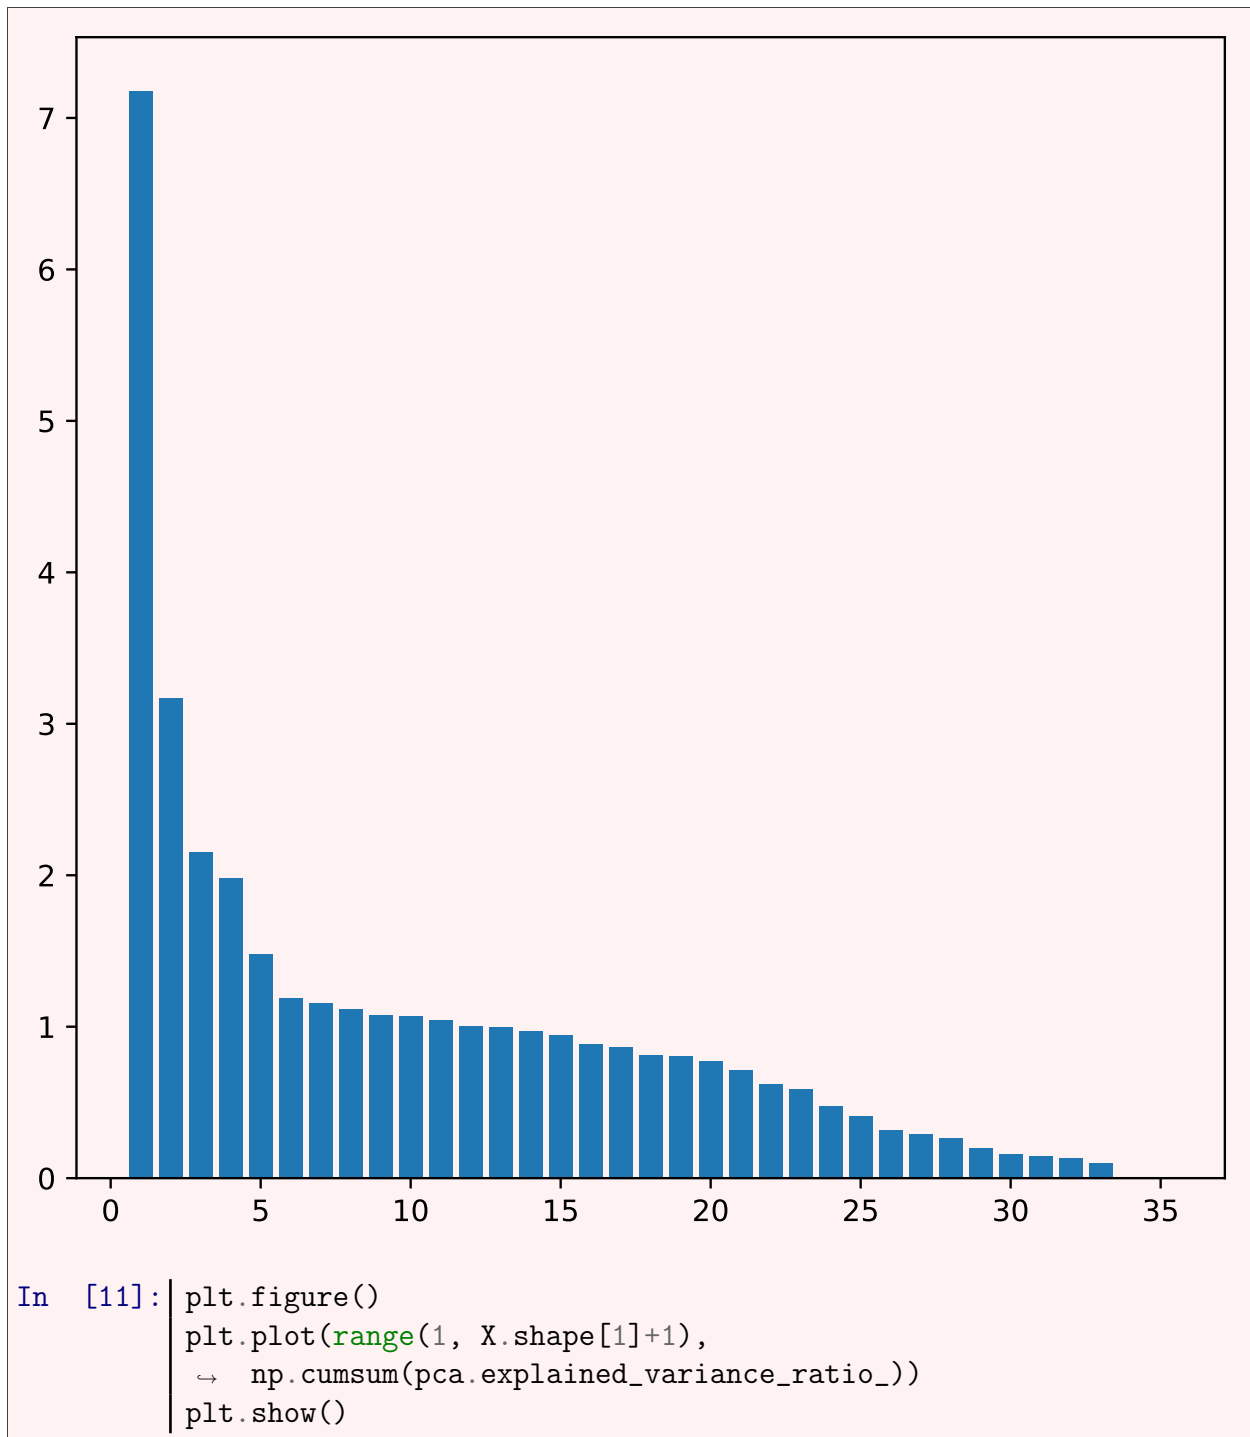
```
In [8]: from sklearn.preprocessing import StandardScaler
        from sklearn.datasets import fetch_openml
        housing = fetch_openml(name="house_prices", as_frame=True,
        → parser="auto")
        X = housing.frame
        X = X.loc[:, ~X.isna().any() &
        → (X.dtypes.astype(str).isin(["float64", "int64"]))]
        X = StandardScaler().fit_transform(X)
        y = housing.target
        pca = PCA()
        pca.fit(X)


Out [8]: PCA()

In [9]: plt.figure()
        plt.bar(range(1, X.shape[1]+1), pca.explained_variance_)

Out [9]: <BarContainer object of 35 artists>

In [10]: plt.show()
```





⑤ Describe the following code. What is it supposed to be doing? Adapt it to determine the optimal number of principal components for the regression task at hand.

```
from sklearn.decomposition import PCA
from sklearn.linear_model import LinearRegression
from sklearn.model_selection import GridSearchCV
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import StandardScaler

pca = PCA()
lin = LinearRegression()
```

```
pca_lin = Pipeline([("scale", StandardScaler()), ("pca", pca), ("lin", lin)])
clf = GridSearchCV(
    estimator=pca_lin,
    cv=10,
    param_grid=dict(pca__n_components=range(1, X.shape[1] + 1)),
)
clf.fit(X, y)
```

This snippet of code is defining a pipeline consisting of a PCA followed by a linear regression on the housing data. The optimal number of components is then computed by cross-validation with 10 folds.

```
In [12]: from sklearn.decomposition import PCA
         from sklearn.linear_model import LinearRegression
         from sklearn.model_selection import GridSearchCV
         from sklearn.pipeline import Pipeline
         from sklearn.preprocessing import StandardScaler

         pca = PCA()
         lin = LinearRegression()
         pca_lin = Pipeline([("scale", StandardScaler()), ("pca", pca),
                               ↪ ("lin", lin)])
         clf = GridSearchCV(
             estimator=pca_lin,
             cv=10,
             param_grid=dict(pca__n_components=range(1, X.shape[1] + 1)),
         )
         clf.fit(X, y)

Out [12]: GridSearchCV(cv=10,
                      estimator=Pipeline(steps=[('scale',
                               ↪ StandardScaler()),
                                                ('pca', PCA()),
                                                ('lin',
                               ↪ LinearRegression())]),
                      param_grid={'pca__n_components': range(1, 36)})

In [13]: print(clf.best_params_)

Out [13]: {'pca__n_components': 33}
```

The best number of principal components using the mean square error and a 10-CV is then ≈ 30 (it might vary)

3 Problem: band reduction in multispectral images

A multispectral image is an image that has several components. For example, a color image has 3 components: red, green and blue and each pixel can be viewed as a vector in \mathbb{R}^3 . More generally a multispectral image of size $N \times M$ with P spectral bands can be stored as a $N \times M \times P$ array. There are $N \times M$ pixels living in \mathbb{R}^P .

When the number of spectral bands P is too large, it is desirable to somehow reduce that number ultimately to 3 for viewing purposes. This process is called band reduction.

Propose a method using the PCA performing a band reduction to 3 bands and use it on the provided multispectral image.

Some multispectral images are available on the internet to test your band reduction algorithm. See for example the following website

- <http://lesun.weebly.com/hyperspectral-data-set.html>

Most of them are available as a Matlab data file (.mat files). It can be loaded with `scipy` with the following function

```
| scipy.io.loadmat
```

You will probably have to reshape arrays. It can be done with the `reshape` method. For example, an array of size $6 \times 6 \times 3$ can be “linearized” using `reshape`

```
| X_lin = X.reshape((-1, 3))
```

the `-1` is automatically inferred from the number of elements in the array. The array is then reshaped into an array of size 36×3 .

It might be handy to be able to rescale the data when it has to belong to some specific range. `scikit-learn` has several rescalers available. For example

```
| from sklearn.preprocessing import MinMaxScaler
```

rescales the data between 0 and 1.

`matplotlib` can display images with the function

```
| plt.imshow
```

Beware of the type of the array (float or integers)!